

One-sided Communication with MPI-2

Rolf Rabenseifner

University of Stuttgart
High-Performance Computing-Center Stuttgart (HLRS)
www.hlrs.de



Acknowledgements

This course is based on the "One-sided" chapter of the MPI-2 tutorial on the MPIDC 2000:


MPI-2: Extensions to the Message Passing Interface

MISSISSIPPI STATE UNIVERSITY¹
**TO PERFORMANCE
IN COMPUTING LAB**
NSF ENGINEERING RESEARCH CENTER



H L R I S²

Anthony Skjellum¹
Purushotham Bangalore¹, Shane Hebert¹
Rolf Rabenseifner²




Synchronization Taxonomy

Message Passing:
explicit transfer, implicit synchronization,
implicit cache operations

Access to other processes' memory:

- **1-sided**
explicit transfer, explicit synchronization,
implicit cache operations (problem!)
- Shared Memory
implicit transfer, explicit synchronization,
implicit cache operations
- shmem interface
explicit transfer, explicit synchronization,
explicit cache operations





MPI-2 One-sided Communication Rolf Rabenseifner
Slide 3 Höchstleistungsrechenzentrum Stuttgart

H L R I S

Cooperative Communication

- MPI-1 supports cooperative or 2-sided communication
- Both sender and receiver processes must participate in the communication





MPI-2 One-sided Communication Rolf Rabenseifner
Slide 4 Höchstleistungsrechenzentrum Stuttgart

H L R I S

One-sided Communication

- Communication parameters for both the sender and receiver are specified by one process (origin)
- User must impose correct ordering of memory accesses




MPI-2 One-sided Communication Rolf Rabenseifner
Slide 5 Höchstleistungsrechenzentrum Stuttgart

H L R I S

One-sided Operations

- Initialization
 - MPI_ALLOC_MEM, MPI_FREE_MEM
 - MPI_WIN_CREATE, MPI_WIN_FREE
- Remote Memory Access (RMA, nonblocking)
 - MPI_PUT
 - MPI_GET
 - MPI_ACCUMULATE
- Synchronization
 - MPI_WIN_FENCE (like a barrier)
 - MPI_WIN_POST / MPI_WIN_START / MPI_WIN_COMPLETE / MPI_WIN_WAIT
 - MPI_WIN_LOCK / MPI_WIN_UNLOCK




MPI-2 One-sided Communication Rolf Rabenseifner
Slide 6 Höchstleistungsrechenzentrum Stuttgart

H L R I S

Window Creation

- Specifies the region in memory (already allocated) that can be accessed by remote processes
- Collective call over all processes in the intracommunicator
- Returns an opaque object of type `MPI_win` which can be used to perform the remote memory access (RMA) operations

```
MPI_WIN_CREATE(base_address, win_size,  
disp_unit, info, comm, win)
```



MPI-2 One-sided Communication Rolf Rabenseifner
Slide 7 Höchstleistungsrechenzentrum Stuttgart


H L R I S

MPI_Put

- Performs an operation equivalent to a send by the origin process and a matching receive by the target process
- The origin process specifies the arguments for both the origin and target process
- The target buffer is at address $\text{target_addr} = \text{win_base} + \text{target_disp} * \text{disp_unit}$

```
MPI_PUT( origin_address, origin_count, origin_datatype,  
target_rank, target_disp, target_count,  
target_datatype, win)
```

Heterogeneous platforms: Use only basic datatypes or derived datatypes without byte-length displacements!



MPI-2 One-sided Communication Rolf Rabenseifner
Slide 8 Höchstleistungsrechenzentrum Stuttgart

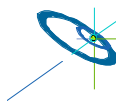
H L R I S

MPI_Get

- Similar to the put operation, except that data is transferred from the target memory to the origin process
- To complete the transfer a synchronization call must be made on the window involved
- The local buffer should not be accessed until the synchronization call is completed

```
MPI_GET( origin_address, origin_count, origin_datatype,  
        target_rank, target_disp, target_count,  
        target_datatype, win)
```

Heterogeneous platforms: Use only basic datatypes or derived datatypes without byte-length displacements!



MPI-2 One-sided Communication Rolf Rabenseifner
Slide 9 Höchstleistungsrechenzentrum Stuttgart

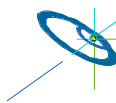
H L R I S

MPI_Accumulate

- Accumulates the contents of the origin buffer to the target area specified using the predefined operation `op`
- User-defined operations cannot be used
- Accumulate is atomic: many accumulates can be done by many origins to one target
-> [may be very expensive]

```
MPI_ACCUMULATE(origin_address, origin_count,  
               origin_datatype, target_rank, target_disp,  
               target_count, target_datatype, op, win)
```

Heterogeneous platforms: Use only basic datatypes or derived datatypes without byte-length displacements!





MPI-2 One-sided Communication Rolf Rabenseifner
Slide 10 Höchstleistungsrechenzentrum Stuttgart

H L R I S


Synchronization Calls

- Active target communication
 - communication paradigm similar to message passing model
 - target process participates only in the synchronization
 - fence or post-start-complete-wait
- Passive target communication
 - communication paradigm closer to shared memory model
 - only the origin process is involved in the communication
 - lock/unlock



MPI_Win_fence

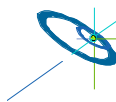

- Synchronizes RMA operations on specified window
- Collective over the window
- Like a barrier
- Should be used before and after calls to put, get, and accumulate
- The `assert` argument is used to provide optimization hints to the implementation
- Used for active target communication



MPI_WIN_FENCE(assert, win)

Start/Complete and Post/Wait, I.

- Used for active target communication with weak synchronization

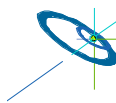
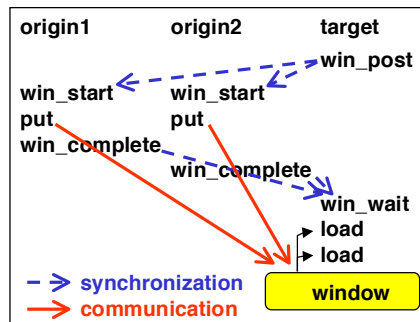


MPI-2 One-sided Communication Rolf Rabenseifner
Slide 13 Höchstleistungsrechenzentrum Stuttgart

HLRS

Start/Complete and Post/Wait, II.

- RMA (put, get, accumulate) are finished
 - locally after win_complete
 - at the target after win_wait
 - local buffer must not be reused before RMA call locally finished
 - communication partners must be known
 - no atomicity for overlapping “puts”
 - assertions may improve efficiency
- > give all information you have

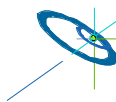
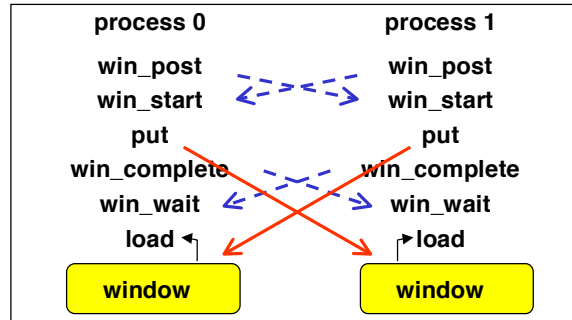


MPI-2 One-sided Communication Rolf Rabenseifner
Slide 14 Höchstleistungsrechenzentrum Stuttgart

HLRS

Start/Complete and Post/Wait, III.

- symmetric communication possible, only win_start and win_wait may block

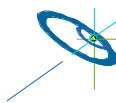



MPI-2 One-sided Communication Rolf Rabenseifner
Slide 15 Höchstleistungsrechenzentrum Stuttgart

HLRS

Lock/Unlock

- Does not guarantee a sequence
- agent may be necessary on systems without (virtual) shared memory
- Portable programs can use lock calls to windows in memory allocated **only** by **MPI_ALLOC_MEM**
- RMA completed after **UNLOCK** at both origin and target



MPI-2 One-sided Communication Rolf Rabenseifner
Slide 16 Höchstleistungsrechenzentrum Stuttgart

HLRS

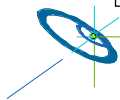
MPI_ALLOC_MEM

MPI_ALLOC_MEM (size, info, *baseptr*)

MPI_FREE_MEM (base)

```

REAL A
POINTER (P, A(100)) ! no memory is allocated
INTEGER (KIND=MPI_ADDRESS_KIND) Size
INTEGER Lng_real, Win, IERR
CALL MPI_TYPE_EXTENT(MPI_REAL, Lng_real, IERR)
Size = 100*Lng_real
CALL MPI_ALLOC_MEM(Size, MPI_INFO_NULL, P, IERR)
CALL MPI_WIN_CREATE(A, Size, Lng_real,
    MPI_INFO_NULL, MPI_COMM_WORLD, Win, IERR)
...
CALL MPI_WIN_FREE(Win, IERR)
CALL MPI_FREE_MEM(A, IERR)
    
```



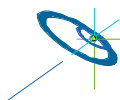
MPI-2 One-sided Communication Rolf Rabenseifner
Slide 17 Höchstleistungsrechenzentrum Stuttgart

H L R I S

Fortran Problems with 1-Sided

Source of Process 1	Source of Process 2	Executed in Process 2
bbbb = 777	buff = 999	register_A := 999
call MPI_WIN_FENCE	call MPI_WIN_FENCE	
call MPI_PUT(bbbb into buff of process 2)		stop application thread
		buff := 777 in PUT handler
		continue application thread
call MPI_WIN_FENCE	call MPI_WIN_FENCE	
	ccc = buff	ccc := register_A

- Fortran register optimization
- Result ccc=999, but expected ccc=777
- How to avoid: (see MPI-2, Chap. 6.7.3)
 - window memory declared in COMMON blocks
i.e. MPI_ALLOC_MEM cannot be used
 - declare window memory as VOLATILE
(non-standard, disables compiler optimization)
 - Calling MPI_Address(buff, idummy_addr, ierror) after 2nd FENCE in process 2

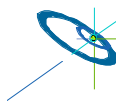


MPI-2 One-sided Communication Rolf Rabenseifner
Slide 18 Höchstleistungsrechenzentrum Stuttgart

H L R I S

One-sided: Summary

- Three one-sided communication primitives provided
 - put / get / ccumulate
- Several synchronization options supported
 - fence / post-start-complete-wait / lock-unlock
- User must ensure that there are no conflicting accesses
- For better performance **assertions** should be used with fence/start/post operations



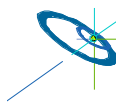
MPI-2 One-sided Communication Rolf Rabenseifner
Slide 19 Höchstleistungsrechenzentrum Stuttgart



MPI-One-sided Exercise 1: Ring communication with fence

- Copy to your local directory:

```
cp ~/MPI/course/C/1sided/ring.c my_1sided_exa1.c  
cp ~/MPI/course/F/1sided/ring.f my_1sided_exa1.f
```
- Tasks:
 - Substitute the non-blocking communication by one-sided communication. Two choices:
 - **either rcv_buf = window**
 - MPI_Win_fence - the rcv_buf can be used to receive data
 - MPI_Put - to write the content of the local variable snd_buf into the remote window (rcv_buf)
 - MPI_Win_fence - the one-sided communication is finished, rcv_buf is filled
 - **or snd_buf = window**
 - MPI_Win_fence - the snd_buf is filled
 - MPI_Get - to read the content of the remote window (snd_buf) into the local variable rcv_buf
 - MPI_Win_fence - the one-sided communication is finished, rcv_buf is filled
 - Compile and run your my_1sided_exa1.c / .f



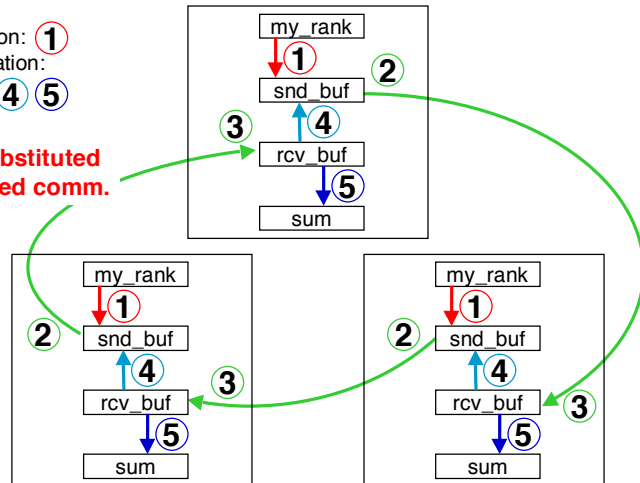
MPI-2 One-sided Communication Rolf Rabenseifner
Slide 20 Höchstleistungsrechenzentrum Stuttgart



ring.c / .f: Rotating information around a ring

Initialization: ①
Each iteration: ② ③ ④ ⑤

to be substituted
by 1-sided comm.



MPI-2 One-sided Communication Rolf Rabenseifner
Slide 21 Höchstleistungsrechenzentrum Stuttgart

H L R I S

MPI-One-sided Exercise 1: additional hints

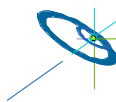
- `MPI_Win_create`:
 - `base` = reference to your `rcv_buf` or `snd_buf` variable
 - `disp_unit` = number of bytes of one int / integer, because this is the datatype of the buffer (=window)
 - `size` = same number of bytes, because buffer size = 1 value
 - `size` and `disp_unit` have different internal representations, therefore:
 - C: `MPI_Win_create(&rcv_buf, sizeof(int), (MPI_Aint) sizeof(int), MPI_INFO_NULL, ..., &win);`
 - Fortran: `INTEGER disp_unit`
`INTEGER (KIND=MPI_ADDRESS_KIND) size`
`CALL MPI_TYPE_EXTENT(MPI_INTEGER, disp_unit, ierror)`
`size = disp_unit * 1`
`CALL MPI_WIN_CREATE(rcv_buf, size, disp_unit, MPI_INFO_NULL, ..., ierror)`
- see MPI-2, page 110

MPI-2 One-sided Communication Rolf Rabenseifner
Slide 22 Höchstleistungsrechenzentrum Stuttgart

H L R I S

MPI-One-sided Exercise 1: additional hints

- MPI_Put or MPI_Get:
 - target_disp
 - C: `MPI_Put(&snd_buf, 1, MPI_INT, right, (MPI_Aint) 0, 1, MPI_INT, win);`
 - Fortran: `INTEGER (KIND=MPI_ADDRESS_KIND) target_disp`
`target_disp = 0`
`CALL MPI_GET(snd_buf, 1, MPI_INTEGER, right, target_disp, 1,`
`MPI_INTEGER, win, ierror)`
- see MPI-2, page 116

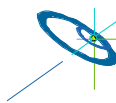


MPI-2 One-sided Communication Rolf Rabenseifner
Slide 23 Höchstleistungsrechenzentrum Stuttgart



MPI-One-sided Exercise 2: Post-start-complete-wait

- Use your result of exercise 1 or copy to your local directory:
`cp ~/MPI/course/C/1sided/ring_1sided.c my_1sided_exa2.c`
`cp ~/MPI/course/F/1sided/ring_1sided.f my_1sided_exa2.f`
- Tasks:
 - Substitute the two calls to `MPI_Win_fence` by calls to `MPI_Win_post` / `_start` / `_complete` / `_wait`
 - Use to group mechanism to address the neighbors:
 - `MPI_COMM_GROUP(comm, group)`
 - `MPI_GROUP_INCL(group, n, ranks, newgroup)`
 - `MPI_COMM_CREATE(comm, group, newcomm)`
 - do not forget **error** with Fortran!
 - Fortran: `integer comm, group, newgroup, newcomm, n, ranks(...)`
 - C: `MPI_Comm comm, newcomm; MPI_Group group, newgroup; int n, ranks[...];`
 - Compile and run your `my_1sided_exa2.c` / `.f`



MPI-2 One-sided Communication Rolf Rabenseifner
Slide 24 Höchstleistungsrechenzentrum Stuttgart

